

Predictive Housing Price Model for Washington State

Group 3

Marriott, Bobby

Martinez, Patricio

Widjaja, Victoria (Tori)

1. Data Importing

Dataset description:

- 21 columns
- 21613 rows

Price Details:

count	21613.00
mean	540088.14
std	367127.20
min	75000.00
25%	321950.00
50%	450000.00
75%	645000.00
max	7700000.00

Name: price, dtype: object

Data Types:

id	int64
date	object
price	float64
bedrooms	float64
bathrooms	float64
sqft_living	float64
sqft_lot	float64
floors	float64
waterfront	int64
view	int64
condition	int64
grade	int64
sqft_above	int64
sqft_basement	int64
yr_built	int64
yr_renovated	int64
zipcode	int64
lat	float64
long	float64
sqft_living15	int64
sqft_lot15	int64

dtype: object

Null Data Counts:

id	0
date	0
price	0
bedrooms	1134
bathrooms	1068
sqft_living	1110
sqft_lot	1044
floors	0
waterfront	0
view	0
condition	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0

dtype: int64

1. Cleaning and Wrangling

Standardizing the date format

```
### working copy ###
date_time_cleaned = df_orig.copy()

### Fixing date format###
date_time_cleaned['date'] = pd.to_datetime(df_orig['date'], format='%Y%m%dT%H%M%S', errors='coerce')

### Save the updated DataFrame to a new CSV file ###
date_time_cleaned.to_csv(os.path.join(log_prefix, "house_file_v2.csv"), index=False)

### sample of the cleaned data ###
display(HTML("<u>Sample of Cleaned Data:</u>"))
print(date_time_cleaned.head())
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	2014-10-13	221900.00	3.00	1.00	1180.00	
1	6414100192	2014-12-09	538000.00	3.00	2.25	2570.00	
2	5631500400	2015-02-25	180000.00	2.00	1.00	770.00	
3	2487200875	2014-12-09	604000.00	4.00	3.00	1960.00	
4	1954400510	2015-02-18	510000.00	3.00	2.00	1680.00	

1. Cleaning and Wrangling

Removing Bedrooms and Bathrooms with a value of 0

```
### Filtering and saving the data ###  
date_time_cleaned = date_time_cleaned[(date_time_cleaned['bedrooms'] != 0) & (date_time_cleaned['bathrooms'] != 0)]  
  
date_time_cleaned.to_csv(os.path.join(log_prefix, 'filtered_date_time_cleaned.csv'), index=False)
```

```
### Checking progress ###  
date_time_cleaned.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	y
0	7129300520	2014-10-13	221900.00	3.00	1.00	1180.00	5650.00	1.00	0	0	...	7	1180	0	1955	
1	6414100192	2014-12-09	538000.00	3.00	2.25	2570.00	7242.00	2.00	0	0	...	7	2170	400	1951	
2	5631500400	2015-02-25	180000.00	2.00	1.00	770.00	10000.00	1.00	0	0	...	6	770	0	1933	

1. Cleaning and Wrangling

Removing the bottom and top 1% of data

```
### Removing the top and bottom 1% of data in the price column ###
working_data = pd.read_csv(os.path.join(log_prefix, 'filtered_date_time_cleaned.csv'))

### Establishing lower and upper bound using the bottom and upper 1% ##
upper_bound = working_data['price'].quantile(0.99)
lower_bound = working_data['price'].quantile(0.01)

display(HTML("<u>Defining Bounds to Remove Outliers:</u>"))
print("Upper bound =", round(upper_bound, 2))
print("\nLower bound =", round(lower_bound, 2))

### Only extracting the data in between the values above the lower bound and values lower than the upper bound ###
working_data2 = working_data[(working_data['price'] >= lower_bound) & (working_data['price'] <= upper_bound)]
```

Defining Bounds to Remove Outliers:

Upper bound = 1965006.6

Lower bound = 154000.0

1. Cleaning and Wrangling

Imputing Part One

```
### Replace NaN values in 'bathrooms' with values from 'avg_bathrooms' ###  
df_clean['bathrooms'] = df_clean['bathrooms'].fillna(df_clean['avg_bathrooms'])
```

```
### Replace NaN values in 'bathrooms' with values from 'avg_bathrooms' ###  
df_clean['bathrooms'] = df_clean['bathrooms'].fillna(df_clean['avg_bathrooms'])
```

```
### Replace NaN values in 'SQFT' with values from 'avg_sqft_living_bin' ###  
df_clean['sqft_living'] = df_clean['sqft_living'].fillna(df_clean['avg_sqft_living'])
```

```
### Replace NaN values in 'sqft_lot' with values from 'avg_sqft_lot' ###  
df_clean['sqft_lot'] = df_clean['sqft_lot'].fillna(df_clean['avg_sqft_lot'])
```

1. Cleaning and Wrangling

Imputing Part Two



```
### Check all Null values have been handled ###  
null_data = df_clean.isnull().sum()  
display(HTML("<u>Null Data Counts:</u>"))  
print(null_data)
```

Null Data Counts:

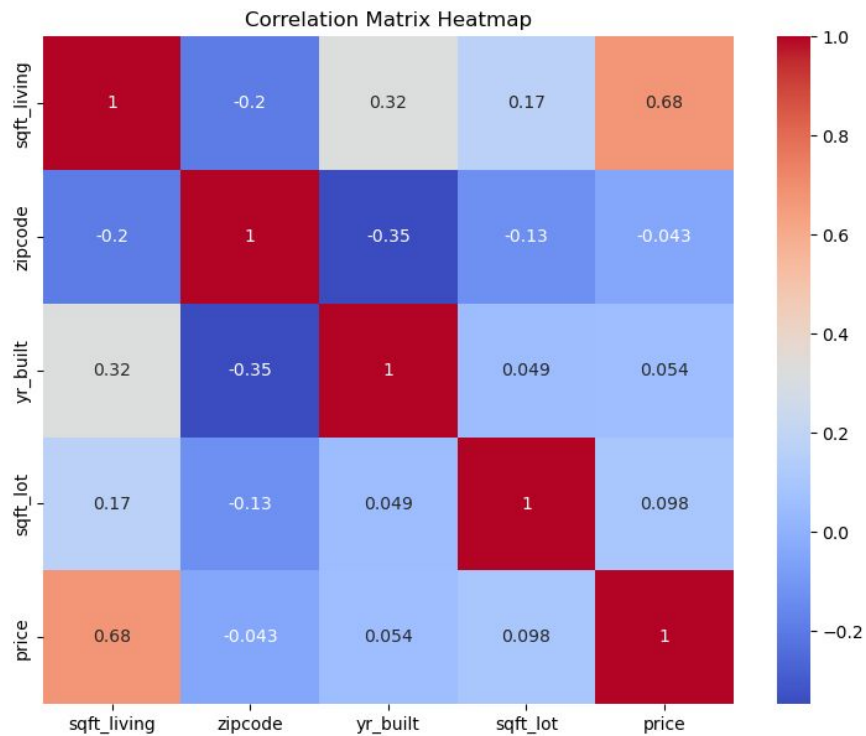
id	0
date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
condition	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0
price_group	0
avg_bedrooms	0
avg_bathrooms	0
avg_sqft_living	0
avg_sqft_lot	0
dtype: int64	

2. Data Analysis and Visualization

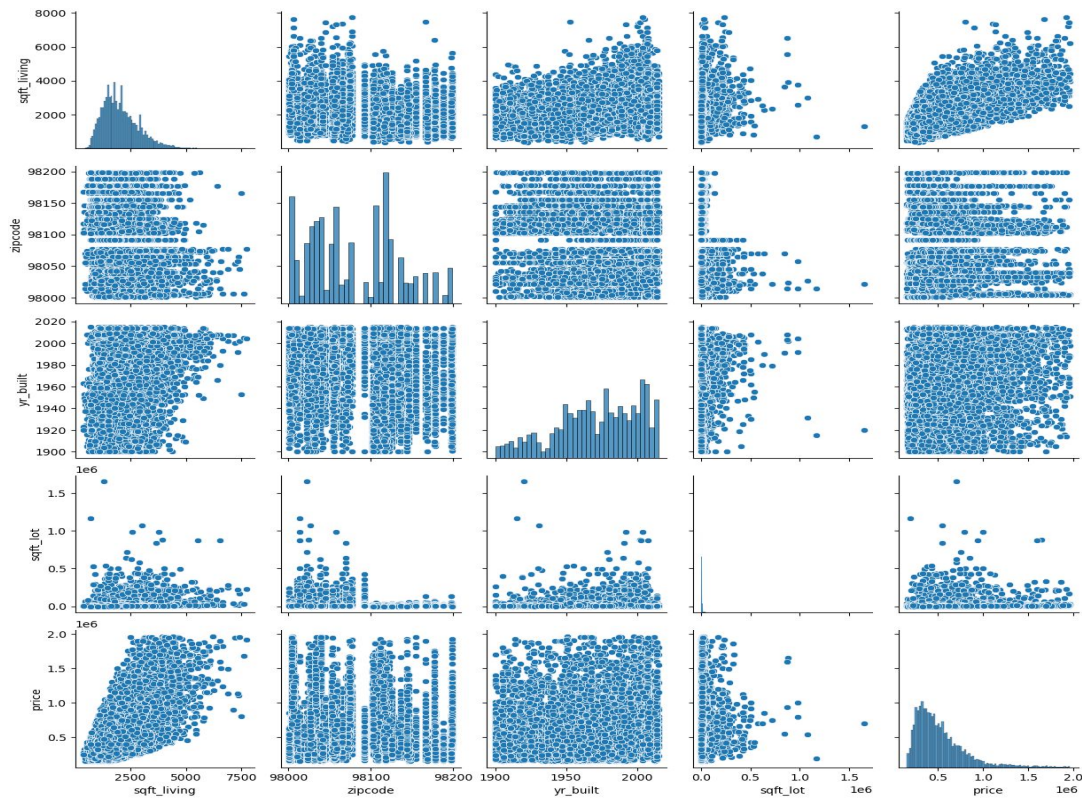
Our goal is to identify key housing features that significantly impact prices in Washington State.



2. Data Analysis and Visualization



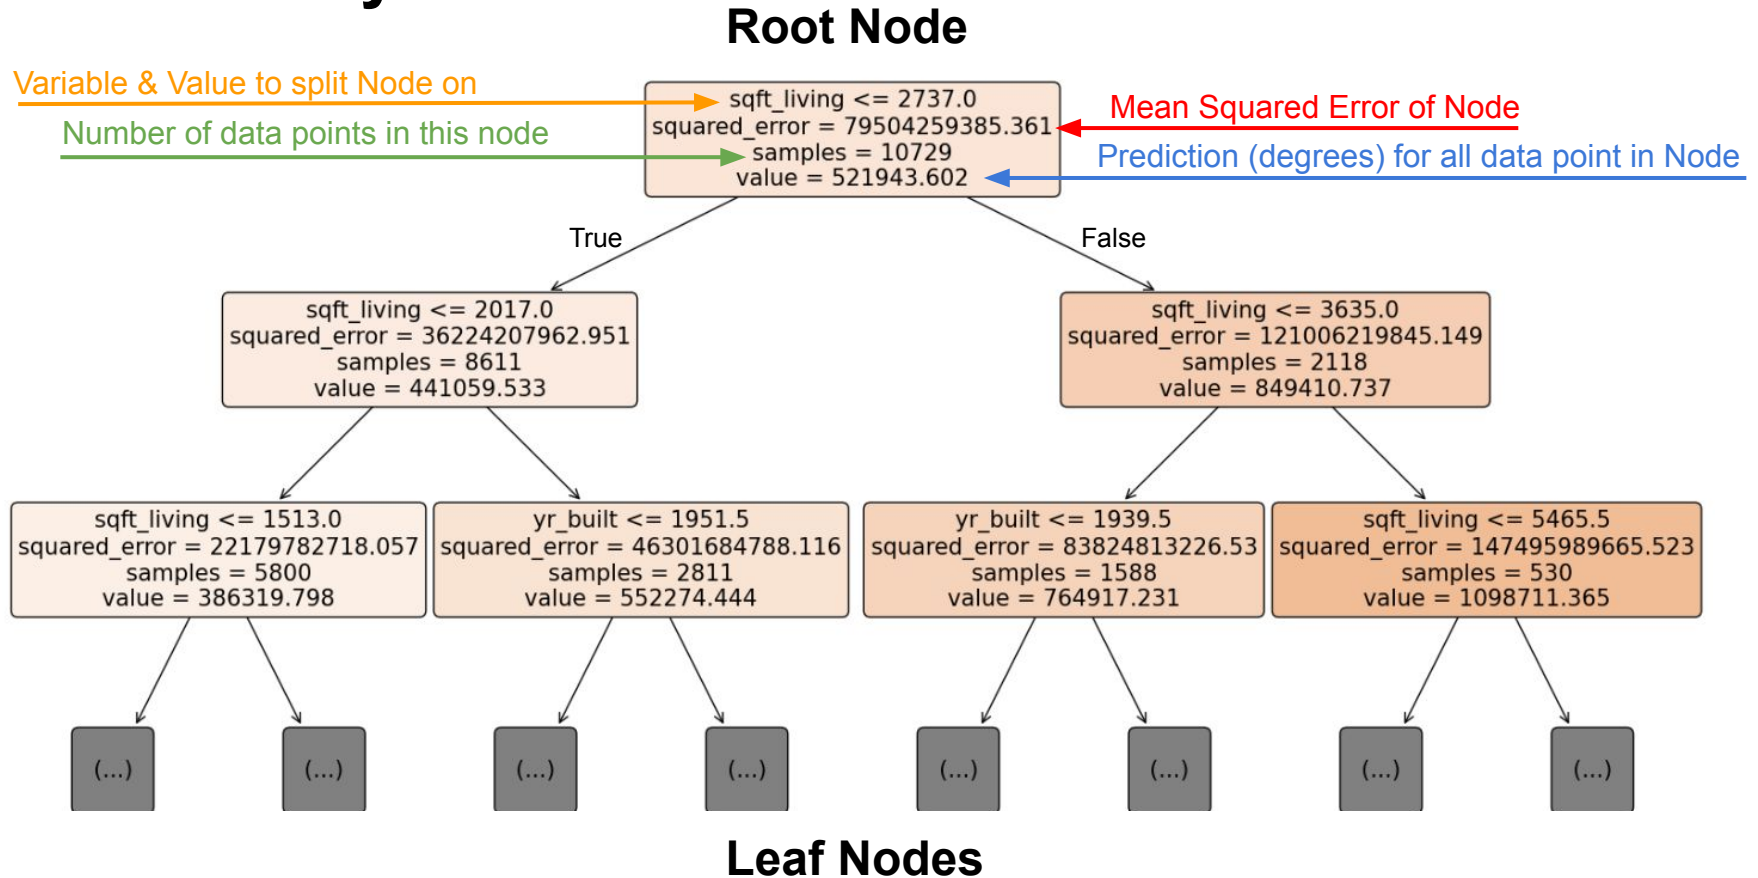
2. Data Analysis and Visualization



3. Data Analytics

- Final Algorithm Type: RandomForestRegressor from sklearn.ensemble
- Dependent Variable: price
- Independent Variables: sqft_living, zipcode, yr_built, sqft_lot

3. Data Analytics



3. Data Analytics

- Mean Absolute Error: \$85,425.55
- Root Mean Squared Error (RMSE): \$135,257.50
- Feature Importances:
 - 'sqft_living' 54.71%
 - 'zipcode' 24.54%
 - 'yr_built' 10.64%
 - 'sqft_lot' 10.11%
- R-squared (R²) Value: 0.7797
- Explained Variance Score: 0.7798

References

Bruce, P., Bruce, A., & Gedeck, P. (2020). Practical Statistics for Data Scientists (Second Edition ed.). O'Reilly Media, Inc.

Koehrsen, W. (2017). Random Forest in Python. Retrieved December 6, 2023 from

<https://towardsdatascience.com/random-forest-in-python-24d0893d51c0>